

Clipping

What happens to objects drawn outside the viewing box?

Let's insert the following square

glBegin(GL_POLYGON)

```
glVertex3f(120.0, 120.0, 0.0);
glVertex3f(180.0, 120.0, 0.0);
glVertex3f(180.0, 180.0, 0.0);
glVertex3f(120.0, 180.0, 0.0);
```

glEnd();

Where did the new square go? OpenGL clips the scene to the viewing box before rendering

Clipping is a stage in the rendering pipeline

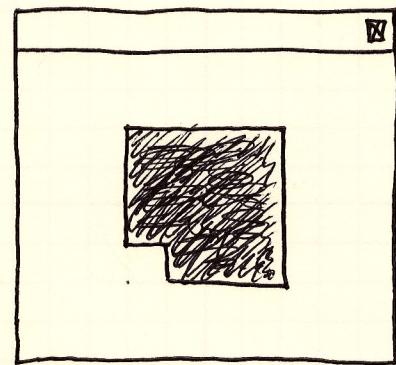
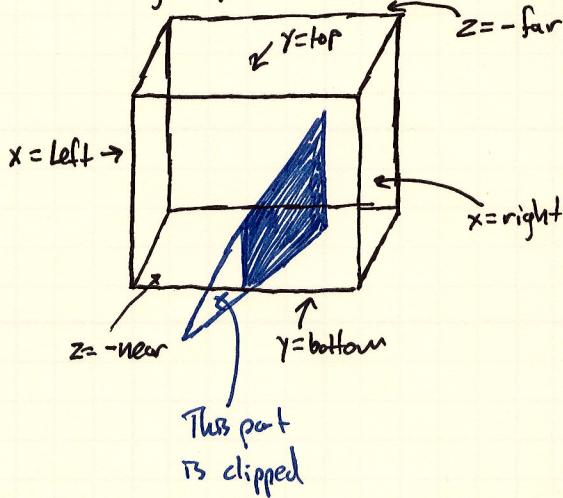
Ex: Redefine the viewing box via glOrtho() such that both squares are visible.

Exp: 1) Replace the original square w/a triangle by commenting out its last vertex

2) Lift the first vertex up the z-axis to 0.5, then further to 1.5, then 2.5 and finally 10.0

The viewing box defines six clipping planes

- Though you can further define additional clipping boundaries



Clipping after raising first vertex

OpenGL Geometric Primitives

- Primitives are used to form the most basic to complex objects in OpenGL
- So far we've seen the **Polygon (GL_POLYGON)** but there are several more
 - **GL_POINTS** - draws a point at each vertex
 - **GL_LINES** - draws a disconnected set of straight line segments between vertices, taken two at a time.
 - **GL_LINE_STRIP** - draws the connected sequence of segments defined by $v_0v_1, v_1v_2, v_2v_3, \dots, v_{n-2}v_{n-1}$ this is called a polyline
 - **GL_LINE_LOOP** - same as **GL_LINE_STRIP**, but includes the additional segment $v_{n-1}v_0$
 - **GL_TRIANGLES** - draws a sequence of triangles using the vertices 3 at a time. If the number of vertices is not a multiple of 3, the remaining 1 or 2 vertices are not used.
 - **GL_TRIANGLE_STRIP** - draws a sequence of triangles, as follows:
 $v_0v_1v_2, v_1v_3v_2, v_2v_3v_4, \dots, v_{n-3}v_{n-2}v_{n-1}$ (if n is odd)
 or
 $v_0v_1v_2, v_1v_3v_2, v_2v_3v_4, \dots, v_{n-3}v_{n-1}v_{n-2}$ (if n is even)
 - **GL_TRIANGLE_FAN** - draws a sequence of triangles around the first vertex, as follows
 $v_0v_1v_2, v_0v_2v_3, \dots, v_0v_{n-2}v_{n-1}$

- NOTE: **GL_POLYGON** has been dropped from the core profile starting in OpenGL 4.0

Drawing Orientation

- The order of vertices determines a polygon's orientation
- Orientation can be either CW or CCW (as seen by the user)
- Orientation is used by OpenGL to determine which face, front or back, the viewer sees.

Polygon Modes

- Drawing mode is set by calling `glPolygonMode(face, mode)`
- face can be: **GL_FRONT**, **GL_BACK** or **GL_FRONT_AND_BACK**
- mode can be: **GL_FILL**, **GL_LINE** or **GL_POINT**
- Whether a polygon is front or back depends on its orientation

OpenGL Geometric Primitives (cont'd)

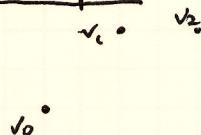
- Point size for GL_POINTS can be set by calling `glPointSize(size)` the default value is 1.0
- Line width for GL_LINES primitives and GL_LINE drawing mode is set by calling `glLineWidth(width)`. The default value is 1.0
- General Definition of a Primitive object

```

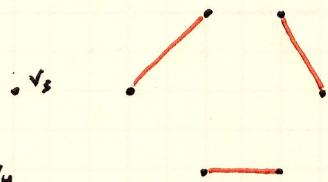
glBegin(primitive type);
    glVertex3f(*, *, *); // v0
    glVertex3f(*, *, *); // v1
    ...
    glVertex3f(*, *, *); // vn-1
glEnd();

```

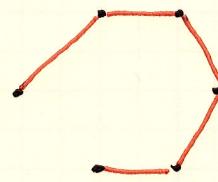
Note: Any configuration commands such as `glLineWidth`, `glPolygonMode`, `glPointSize` must be called outside of `glBegin()` ... `glEnd()` and before the block they pertain to.

Examples

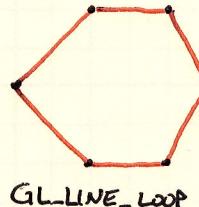
GL_POINTS



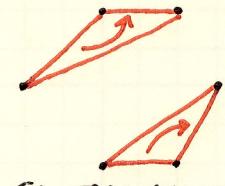
GL-LINES



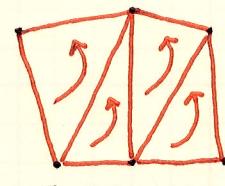
GL-LINE-STRIP



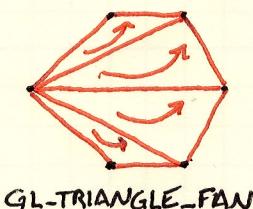
GL-LINE_LOOP



GL_TRIANGLES



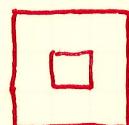
GL-TRIANGLE-STRIP



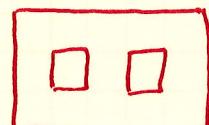
GL-TRIANGLE_FAN

Ex: Replace the GL_POLYGON primitive in `square.cpp` with each of the above. What happens?

Ex: Using only ~~2~~ 2 GL-TRIANGLE-FAN's draw a square annulus.



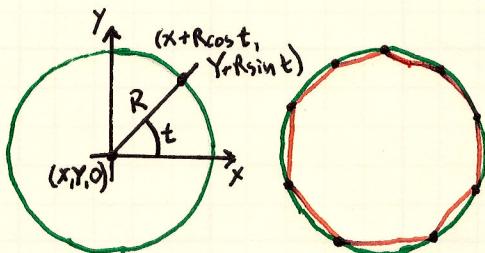
Then draw this:



Approximating Curved Objects

Q: If all of our primitives are composed of points, straight lines and flat segments, then how can we draw curved objects such as circles, ellipses, or beer cans?

- Let us begin by considering a circle ([download + review circle.cpp](#))



- We construct a circle by sampling points using its parametric equations:

$$\begin{aligned} x &= X + R \cos t \\ y &= Y + R \sin t \quad 0 \leq t \leq 2\pi \end{aligned}$$

- Implementation:

```

glBegin(GL_LINE_LOOP);
for (i = 0; i < numVertices; i++)
{
    glColor3f((float) rand() / (float) RAND_MAX,
              (float) rand() / (float) RAND_MAX,
              (float) rand() / (float) RAND_MAX);
    glVertex3f(X + R * cos(t), Y + R * sin(t), 0.0);
    t += 2 * PI / numVertices;
}
glEnd();
    
```

- The larger the ~~number~~ of samples the more precise of a circle that ~~can~~ can be drawn

◦ EXP: Modify this to construct a flat 3-turn spiral

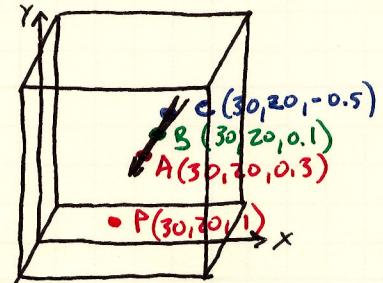


◦ EXP: Modify this to construct a circular annulus



The Depth Buffer

- The depth buffer (or z-buffer) allows OpenGL to eliminate obscured parts of objects prior to rendering
- That is, a point of an object is not drawn if its projection is hidden by another object.
- This process goes by many names
 - Hidden Surface Removal
 - Depth Testing
 - Visibility Determination
- For Orthographic projection this works as follows.



1. We fix an X, Y value as X, Y from the viewing box
2. $S \leftarrow$ set of all points of objects in viewing box along line $L = (X, Y, 0)$
3. S is of form $S = \{(X, Y, z)\}$, where Z varies depending on points
4. $Z \leftarrow \arg \max_z S$
5. $P \leftarrow \text{rgb}(X, Y, Z)$ // all points in S project to $P = (X, Y, -\text{near})$, thus

~~the OpenGL~~

- In the GPU the z-buffer is a block of memory containing z-values (one per pixel)
 - As primitives are processed their z-values are compared against the one currently in the buffer.
 - If the incoming z-value is greater it replaces the current one and the rgb values of the new pixel replace the current values

Note: The GPU uses a scaled value between 0 and 1 where 0 is the near face and 1 is the far face. But, in OpenGL we use world space, which reverses this relationship

- To enable this add the following code:

- `glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT)` in `drawScene` to clear the depth buffer
- `glEnable(GL_DEPTH_TEST)` in `drawScene()` to enable HSR
use `glDisable()` to disable it
- `glutInitDisplayMode(... | GLUT_DEPTH)` in `main()` will initialize the depth buffer.

Ex: Make a bullseye with increasing diameter disks of differing colors. Use at least 5 disks

Perspective Projection

- Motivating Example: a helix governed by the following equations

$$x = R \cos t$$

$$y = R \sin t \quad \text{for } -10\pi \leq t \leq 10\pi$$

$$z = t - 60$$

- Download and review `helix.cpp` from moodle
- Run `helix.cpp`, what do you see?

- Issue: We apparently see only a circle

- Perhaps we can fix this by: replacing `glVertex3f(R * cos(t), R * sin(t), t - 60.0)` with `glVertex3f(R * cos(t), t, R * sin(t), -60.0)`

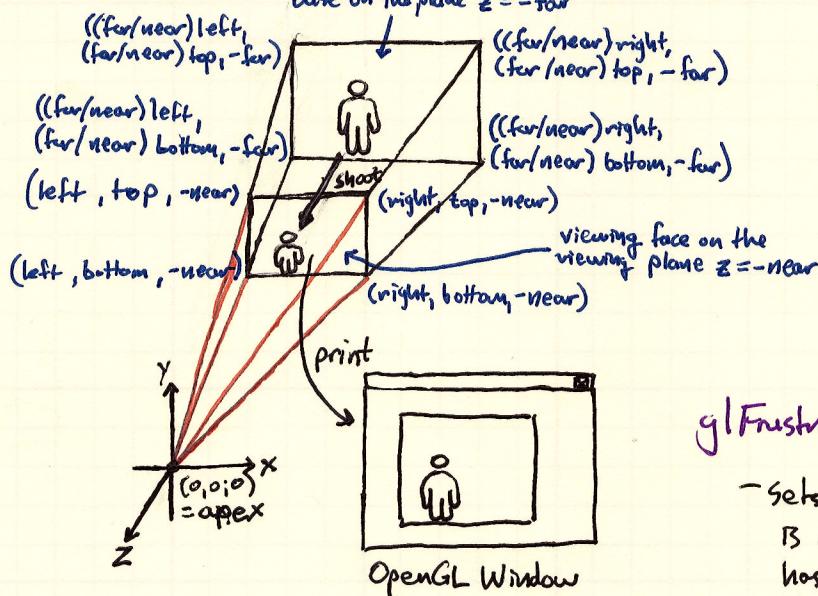
- But this still doesn't seem to do much better!

- Why are these problems occurring?

- Orthographic projection squashes a dimension \Rightarrow Bad for 3D scenes

- Solution: Perspective projection via `glFrustum()`

base on the plane $z = -\text{far}$



Note: We can think of the apex as the location of a point camera

`glFrustum(left, right, bottom, top, near, far)`

- Sets up the viewing frustum which is a truncated pyramid whose top has been cut off by a plane parallel to its base.

- Values of `glFrustum()` are typically set

- to be symmetrical about the z-axis

- right and top are positive

- left and bottom are their negatives

- near and far should both be positive and $\text{near} < \text{far}$

Ex: Replace the `glOrtho(...)` call in `helix.cpp` with:

`glFrustum(-5.0, 5.0, -5.0, 5.0, 5.0, 100.0)`

Perspective vs. Orthographic

- Both use a two-step "print and shoot" rendering process
- Perspective is no longer perpendicular, rather each point is projected along the line joining it to the apex.
- Perspective causes **foreshortening** because objects further from the apex appear smaller
 - This is called ~~is~~ perspective transformation
- The second step is exactly the same for both modes
- Perspective is closer to reality, in that it mimics how images are formed on the retina
- ~~both can have~~
- Both orthographic and perspective projections ~~converge~~ converge to a **center of projection (COP)**
 - For perspective, this is a regular point with finite coordinates
 - For orthographic, this is a point infinitely far away and thus the lines are parallel.
- Both of these projections allow OpenGL to implement the **synthetic-camera model**

Ex: Try the following glFrustum settings in helix.cpp

glFrustum(-5.0, 5.0, -5.0, 5.0, 5.0, 120.0)

glFrustum(-5.0, 5.0, -5.0, 5.0, 10.0, 100.0)

glFrustum(-5.0, 5.0, -5.0, 5.0, 2.5, 100.0)

glFrustum(-10.0, 10.0, -10.0, 10.0, 5.0, 100.0)

What are the effects?

Def: Take a moment and work out an algorithm for HSR in perspective projection.

OpenGL Program Walkthrough (square.cpp)

- main()

- glutInit(...); - initializes the FreeGLUT Library
- glutInitContextVersion(...);
- glutInitContextProfile(GLUT_COMPATIBILITY_PROFILE);

sets the OpenGL context used by FreeGLUT and allows for backwards compatibility

- glutInitDisplayMode(GLUT_SINGLE | GLUT_RGBA); - single-buffered frame with each pixel having red, blue, green, and alpha values.
- glutInitWindowSize(500, 500); } set initial window size and position conditions
- glutInitWindowPosition(100, 100); }
- glutCreateWindow(...); - actually create the OpenGL context and its window with the supplied title
- glutDisplayFunc(drawScene); } register callback routines
- glutReshapeFunc(resize); }
- glutKeyboardFunc(keyInput); }
- glewExperimental = GL_TRUE; } initializes GLEW to load extensions
- glewInit(); }
- setup() - invoke initialization function
- glutMainLoop() - begin event-processing loop, calling registered callbacks as needed

- drawScene() - already covered in-depth

- resize(int w, int h) - handles when the window is resized

- glViewport(0, 0, w, h) - specifies the portion of the window in which rendering will take place.

- glMatrixMode(GL_PROJECTION); } activate the projection matrix stack
- glLoadIdentity(); } place identity matrix on top of stack
- glOrtho(...); } multiply identity matrix by result of glOrtho command

- glMatrixMode(GL_MODELVIEW); } activate modelview matrix stack
- glLoadIdentity(); } place identity matrix on top of this stack

- keyInput(unsigned char key, int x, int y)

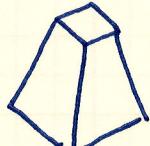
- when a key is pressed its ASCII value is passed to the char parameter along with the current value of the x, y mouse location

For Next Time

- Review Lectures 02 and 03
- Read Chapter 3 Sections 3.1 – 3.7
- Review Lecture 04 when it becomes available
- Begin working on Assignment 01
- Come to class!

Additional Notes

- Can you create these shapes?



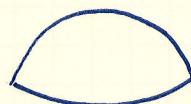
Square lamp
shade



Conical
lamp shade



Rugby
ball



Hemisphere